Lakkaraju Anjaneya Vara Prasad

lakkaraju-anjaneya-vara-prasad-00070b235 |
anjaneyavaraprasad01@gmail.com | 9573294994

Azure Data Engineer | ADF - Databricks (PySpark) - Power BI - Data Quality & Optimization

Summary

Highly motivated Azure Data Engineer with 2 years of experience designing and developing scalable cloud-based data pipelines using Azure Data Factory, Databricks (PySpark), ADLS Gen2, and Synapse Analytics. Proven ability to automate ETL processes, optimize PySpark workloads, and ensure high data quality for analytics and reporting in banking domains. Microsoft Certified Azure Data Engineer Associate (DP-203).

Work Experience

Programmer Analyst – Azure Data Engineer Dec 2023 – Present

- Designed and automated end-to-end ETL pipelines ingesting 1–2 million banking transactions/day from SQL Server, REST API data sources, and flat files ingested from NAS (Network Attached Storage) paths into Azure Data Lake Storage Gen2 (ADLS Gen2) and Azure Synapse Analytics using Azure Data Factory (ADF) and Databricks (PySpark).
- Improved PySpark job performance by 15–20% through optimization strategies such as partitioning, caching, and optimized joins, enhancing workload efficiency and reducing overall compute costs.
- Developed and maintained 4–7 production-grade ETL pipelines with automated error handling and scheduling, reducing manual monitoring by 20% and lowering pipeline failure incidents by 15%.
- Authored advanced SQL validation queries that improved reporting data accuracy by 10%, ensuring reliable datasets consumed by Power BI dashboards for fraud detection and customer analytics.
- Collaborated with business analysts, QA engineers, and senior developers to deliver compliant, high-quality data solutions in an Agile/Scrum environment.

Projects

Streaming state management for user session tracking using pyspark:

- This project leverages PySpark Structured Streaming to implement real-time session tracking and stateful analytics.
- It processes a continuous stream of employee activity events—such as logins, clicks, and logouts—to dynamically manage session states.
- By applying techniques like sessionization, stateful aggregation, and windowed operations, the system provides live insights into session durations, user engagement, and cumulative actions.
- The project showcases the power of streaming state management in operational analytics, enabling responsive dashboards and real-time decision-making across domains like HR, security, and product usage.

Cloud Data Orchestration: From Buckets To BigQuery With Dataflow and Cloud functions:

- This project automates the end-to-end pipeline for ingesting, processing, and loading data from Google Cloud Storage into BigQuery using Cloud Functions and Dataflow.
- The workflow is event-driven: when a new file is uploaded to a bucket, a Cloud Function is triggered to launch a Dataflow job that transforms and loads the data into BigQuery.
- The solution demonstrates scalable and efficient cloud-native orchestration, highlighting seamless integration across Google Cloud Platform (GCP) services
- It emphasizes real-time responsiveness, modular architecture, and cost-effective data movement.

Education

2019 - 2023	Bachelor's Degree at **Vasireddy Venkatadri Institute of Technology**	(GPA: 8.0/10.0)
2017 - 2019	Class 12th Board of Intermediate Education	(GPA: 9.04/10.0)
2016 - 2017	Class 10th Board of Secondary Education	(GPA: 9.8/10.0)

Certifications

- Microsoft Certified: Azure Data Engineer Associate (DP-203)
  Provider: Microsoft
  Completion ID: C1164443AF87DC47
  Valid until: March 2026
  View Credential

Skills

Technical Skills

- Cloud Platforms: AWS (Glue, EMR, Redshift, Lambda, S3), Azure (ADF, Synapse), GCP (BigQuery, Dataflow)
- Big Data: Spark (PySpark), Hadoop, Hive, Sqoop, HDFS
- Azure Databricks, ADLS Gen2, Azure Synapse, Azure SQL, MySQL
- ETL automation, SQL (Advanced), Python
- Power BI dashboards, data validation, governance

Soft Skills

- Communication
- Teamwork
- Problem-solving
- Attention to detail
- Time management