



Lakkaraju Anjaneya Vara Prasad

Programmer Analyst

✉ anjaneyavaraprasad01@gmail.com ☎ 9573294994 📍 hyderabad,India

🌐 <https://www.linkedin.com/in/lakkaraju-anjaneya-vara-prasad-00070b235> 📅 04/08/2001 🇮🇳 Indian

👤 **Headline**

Azure Data Engineer | Azure Data Factory - Databricks (PySpark) - Power BI - Data Quality & Optimization

👤 **Profile**

Data Engineer with 2+ years of experience building scalable cloud-based data pipelines using Azure Data Factory, Databricks (PySpark), ADLS Gen2, and Synapse Analytics. Skilled in automating ETL processes, optimizing Spark workloads, and ensuring data quality for analytics in banking. Microsoft Certified: Azure Data Engineer Associate (DP-203), with experience in Agile/Scrum, incident resolution, and performance tuning.

🎓 **Education**

Bachelor of Technology (B.Tech), Vasireddy Venkatadri Institute of Technology 06/2019 – 05/2023
CGPA : 8.0 Guntur, India

💼 **Professional Experience**

Programmer Analyst, Cognizant Technology Solutions 07/2024 – Present
hyderabad, India

- Designed and automated end-to-end ETL pipelines ingesting 1–2 million banking transactions/day from SQL Server, REST API data sources, and flat files ingested from NAS (Network Attached Storage) paths into Azure Data Lake Storage Gen2 (ADLS Gen2) and Azure Synapse Analytics using Azure Data Factory (ADF) and Databricks (PySpark).
- Improved PySpark job performance by 15–20% through optimization strategies such as partitioning, caching, and optimized joins, enhancing workload efficiency and reducing overall compute costs.
- Developed and maintained 4–7 production-grade ETL pipelines with automated error handling and scheduling, reducing manual monitoring by 20% and lowering pipeline failure incidents by 15%.
- Authored Structured Query Language (SQL) validation queries that improved reporting data accuracy by 10%, ensuring reliable datasets consumed by Power BI (data validation, governance, reporting enablement) dashboards for fraud detection and customer analytics.

– Collaborated with business analysts, QA engineers, and senior developers to deliver compliant, highquality data solutions in an Agile/Scrum environment.

Skills

Technical Skills:

- Cloud Platforms:
- Amazon Web Services (AWS): Glue, EMR, Redshift, Lambda, S3
- Microsoft Azure: Azure Data Factory (ADF), Synapse Analytics, Databricks (PySpark), Azure Data Lake Storage Gen2 (ADLS Gen2), Azure SQL
- Google Cloud Platform (GCP): BigQuery, Dataflow
- Big Data: Apache Spark (PySpark), Hadoop, Hive, Sqoop, HDFS
- Databases: SQL Server, Azure SQL
- ETL Automation: Designed and scheduled pipelines using Azure Data Factory (ADF) and Databricks (PySpark), with Python scripting
- Data Quality: Validation and governance for reporting pipelines

Soft Skills:

- Fast learner with adaptability
- Strong work ethic
- Communication
- Teamwork
- Problem-solving
- Attention to detail
- Time management

Certificates

- Microsoft Certified: Azure Data Engineer Associate (DP-203) 

Projects

Streaming state management for user session tracking using pyspark:

- This project leverages PySpark Structured Streaming to implement real-time session tracking and stateful analytics.
- It processes a continuous stream of employee activity events—such as logins, clicks, and logouts—to dynamically manage session states.
- By applying techniques like sessionization, stateful aggregation, and windowed operations, the system provides live insights into session durations, user engagement, and cumulative actions.
- The project showcases the power of streaming state management in operational analytics, enabling responsive dashboards and real-time decision-making across domains like HR, security, and product usage.

Cloud Data Orchestration: From Buckets To BigQuery With Dataflow and Cloud functions

- This project automates the end-to-end pipeline for ingesting, processing, and loading data from Google Cloud Storage into BigQuery using Cloud Functions and Dataflow.
- The workflow is event-driven: when a new file is uploaded to a bucket, a Cloud Function is triggered to launch a Dataflow job that transforms and loads the data into BigQuery.
- The solution demonstrates scalable and efficient cloud-native orchestration, highlighting seamless integration across Google Cloud Platform (GCP) services
- It emphasizes real-time responsiveness, modular architecture, and cost-effective data movement.